



Drupal AI Factchecker

Claim-level verification on your own corpus.

powered by **AMD**.

Team: Drupal AI Router (Leo Fishman)

License: GPL-2.0+, published on drupal.org

The Problem with Current Tools

- 🎮 Pre-existing software is **excessively expensive** and operates as a closed, proprietary system.
- 🔒 Universities and governments have **almost zero control** over how their sensitive data is processed.
- ✗ Tools detect **copying**, not falsehoods. AI-generated essays are "original" and undetectable.
- ⚠️ Editors and educators lack an in-workflow tool to verify factual accuracy.



The Solution: Native Integration & Control



Where the content lives, under your control.

Institutions facing this problem (governments, universities) already run Drupal massively (.edu, .gov).

Unlike closed SaaS products, a **highly configurable open-source** Drupal module offers **tight integration** with your own content, keeping your data strictly in-house.

The Product: Native Integrity Suite




1. Factcheck

Atomic claim extraction and evidence retrieval. Per-claim verdict (SUPPORTED / UNSUPPORTED / CONTRADICTED) with auditable score.



2. AI Likelihood

0–100 score with rationale. Honestly framed: presented as a contextual hint, not a definitive mathematical proof.



3. Verbatim Plagiarism

Identifies the text's most distinctive sentences and searches them as exact phrases on the web to detect unoriginal copies.

Key Differentiators



Evidence Cascade

The university's OWN indexed corpus is prioritized first. Verdicts are securely grounded in your institutional knowledge.



Distrust List

Claims echoed only by known misinformation domains are marked as tainted. Misinformation sites asserting a claim counts AGAINST it.



100% Auditable

Every model call and routing decision is meticulously logged. No other tool combines local evidence with transparent logging.


Powered by AMD Instinct & Ryzen

All inference (claim extraction + verification) runs on an **AMD Instinct GPU**: ROCm 7.2 + vLLM 0.16 serving Gemma.

Student work NEVER leaves institution-controlled hardware.

Local GPU Execution

rocm-smi execution inhouse



Drupal AI Integration

AMD server endpoint in Drupal

Edit AMD Instinct GPU (ROCm + vLLM)

Server name *

AMD Instinct GPU (ROCm + vLLM) Machine name: amd_vllm

A human-readable name for this server (e.g. "Ollama Local", "GPU Chat Server").

Connection

Backend *

OpenAI-compatible

The protocol this server speaks: OpenAI-compatible covers llama.cpp, Ollama, vLLM, LM Studio and any similar local or remote server. Hosted services (OpenRouter, Hugging Face, Fireworks, Ollama Cloud, LiteLLM/amazonai) have dedicated backends with better model support.

Host Name

http://ollama

The host name including protocol. Local examples: http://127.0.0.1 (llama.cpp, vLLM, LM Studio, LiteLLM), http://host.docker.internal (from Docker), Remote example: https://api.openai.com.

Port

11434

Port number. Leave empty for default. Common defaults: Ollama 11434, llama.cpp 8080, vLLM 8000, LM Studio 1234, LiteLLM 4000. Remote HTTPS APIs usually need no port (443).

API Key

- Select a key -

Required for hosted services. Optional for local servers; leave empty when unauthenticated (llama.cpp, Ollama, LM Studio), set for vLLM or LiteLLM with an api-key configured. Missing key? [Create one in a new tab](#), then press Refresh keys.

Refresh keys

Timeout (seconds)

120


Request timeout in seconds. Increase for slow models or large inputs.

Test connection & list models

Architecture Flow


- **Trigger:** Drupal Node → Content Scan
- **Routing:** ai_provider_universal router
- **Processing:** AMD GPU (vLLM/Gemma)
- **Retrieval:** Vector index (local corpus) + Tavily (web)
- **Auditing:** Securely logged decisions report

Everything—models, backends, evidence sources, trust lists—is Drupal configuration, not code.



Real and Ready

-  Built as a real, production-ready open-source module on **drupal.org**.
-  Built directly on top of the official Drupal AI ecosystem.
-  Backed by months of prior foundational work.
-  **Hackathon Delta:** Adding the critical AMD serving layer and defining the university use case.



Market Opportunity

Sizing the content integrity and plagiarism detection landscape.

\$1.2B

PLAGIARISM DETECTION MARKET

Ready for Enterprise Budgets

Universities already allocate massive budgets for content integrity tooling. This solution slots perfectly into their existing CMS at a fraction of the cost.

Multi-Sector Expansion

Beyond academia, the exact same pipeline serves the urgent needs of newsrooms, investigation academies, and government content teams which most of them are already in Drupal.

Claim-level truth.

Where the content lives, **on AMD.**



GitHub Repo & Demos

<https://github.com/leofishman/amd-hackathon>

<https://amd-hackathon.fishman.work/>



Demo Video Link

<https://www.youtube.com/watch?v=WqwtFhDbv20>



drupal.org Project Page

https://www.drupal.org/project/ai_provider_universal